Scheduler Performance Comparison

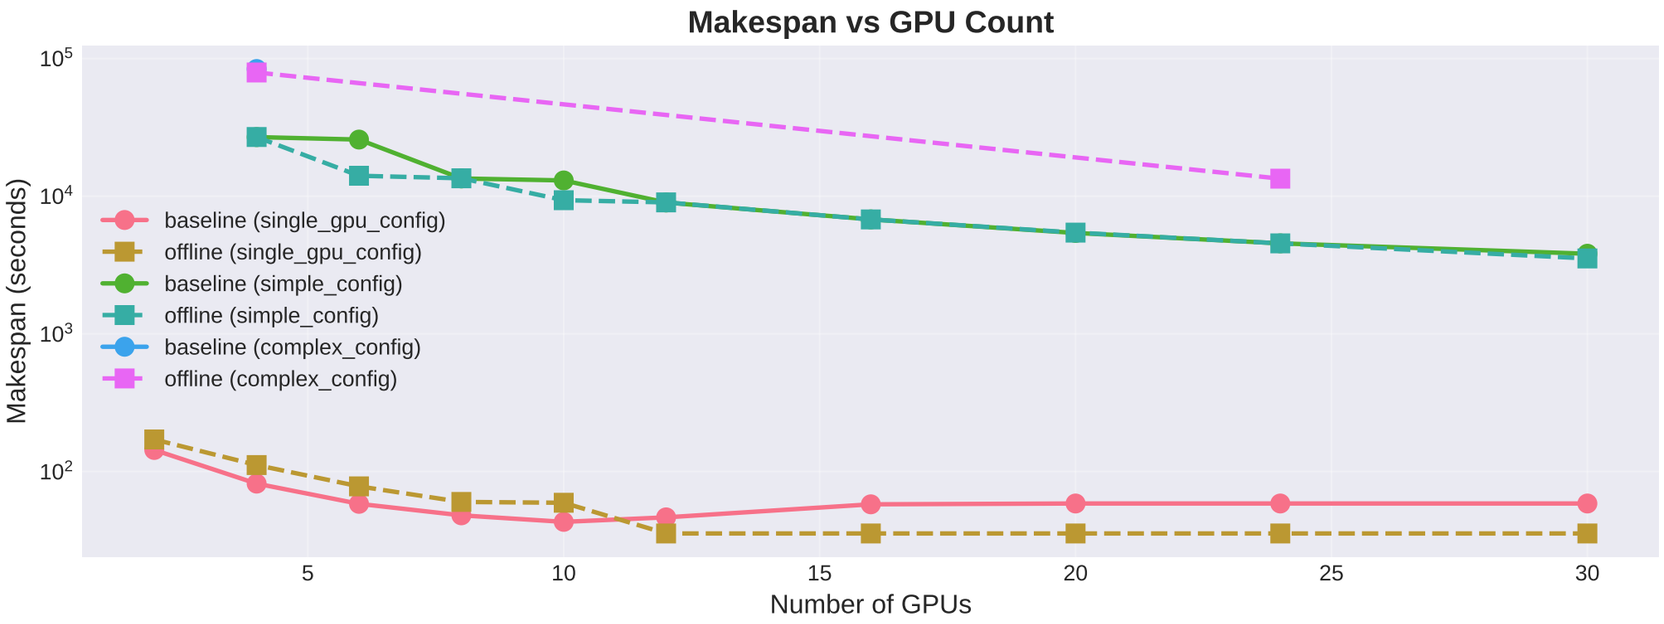

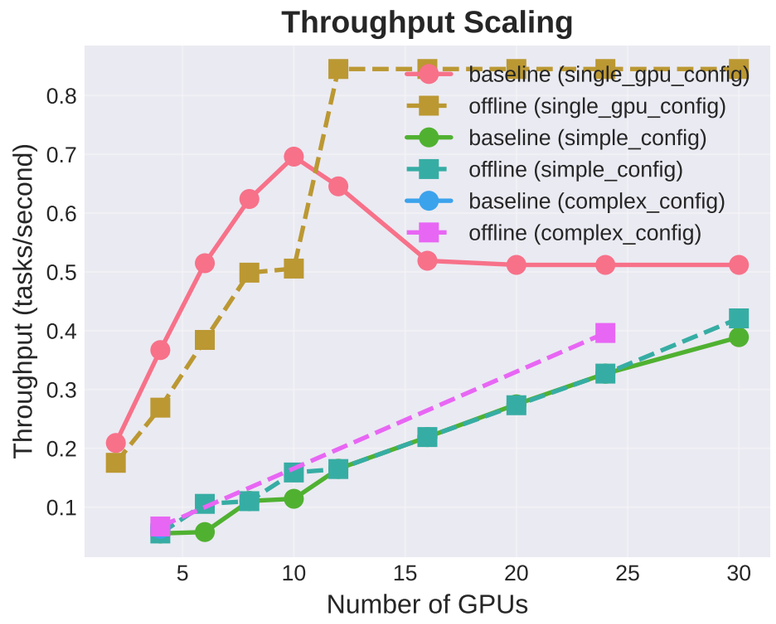

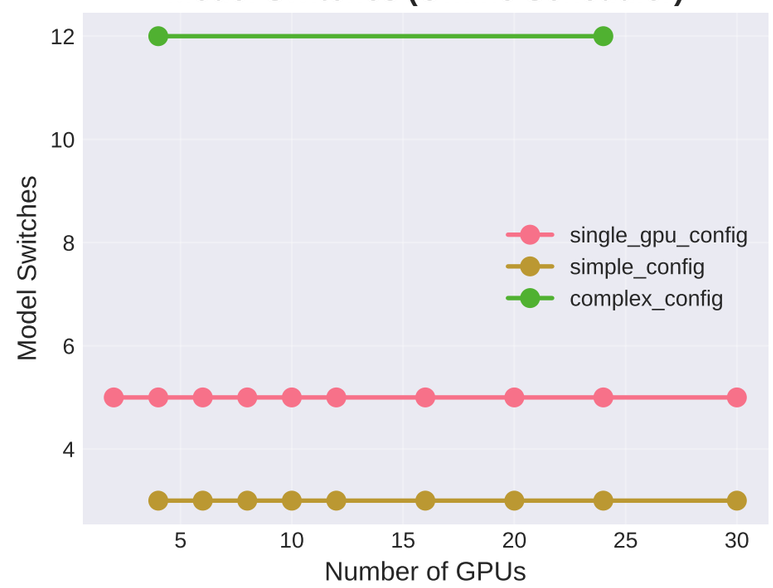

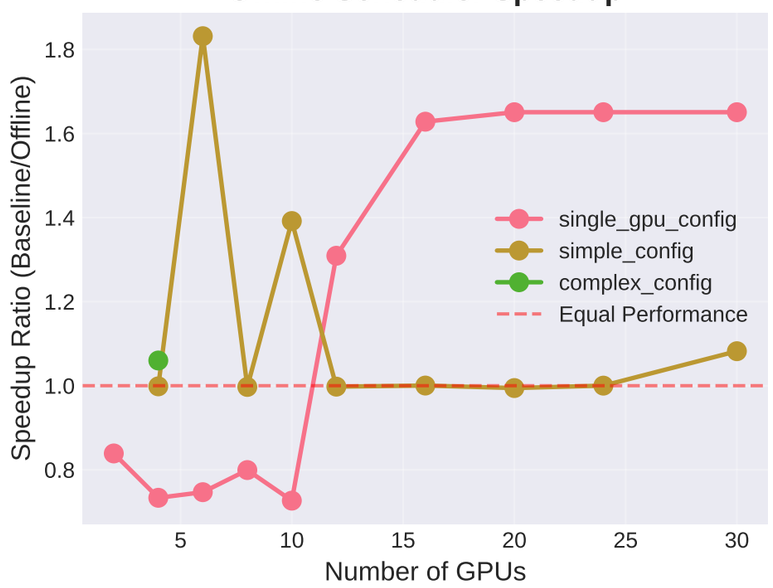

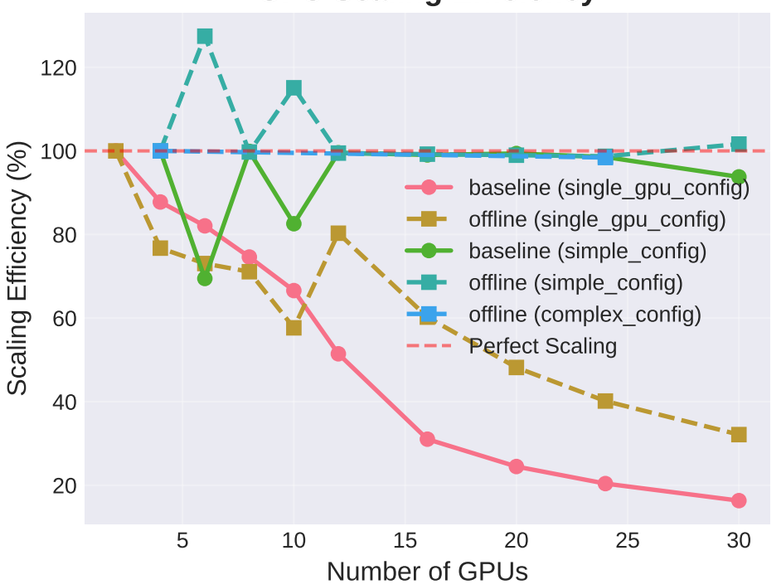


```
Performance Summary:

SINGLE_GPU_CONFIG:
  baseline: Avg makespan=65.5s, Avg throughput=0.511 tasks/s
         Best performance at 10 GPUs (43.1s)
  offline: Avg makespan=65.8s, Avg throughput=0.606 tasks/s
         Best performance at 12 GPUs (35.5s)

SIMPLE_CONFIG:
  baseline: Avg makespan=12072.9s, Avg throughput=0.190 tasks/s
         Best performance at 30 GPUs (3819.5s)
  offline: Avg makespan=10347.5s, Avg throughput=0.204 tasks/s
         Best performance at 30 GPUs (3529.8s)

COMPLEX_CONFIG:
  baseline: Avg makespan=83900.0s, Avg throughput=0.063 tasks/s
         Best performance at 4 GPUs (83900.0s)
  offline: Avg makespan=46280.0s, Avg throughput=0.232 tasks/s
         Best performance at 24 GPUs (13405.2s)

Speedup Analysis:
  single_gpu_config: Offline is 1.00x faster on average
  simple_config: Offline is 1.17x faster on average
  complex_config: Offline is 1.81x faster on average
```